A simulation-based approach for quantifying and partitioning uncertainty to improve forecasts of dynamic processes

Andrew T. Tredennick[1], Peter B. Adler[1], and others[2]

[1] Department of Wildland Resources and the Ecology Center, Utah State University, Logan, UT, United States

[2] Department of somewhere

Abstract

Making informed decisions in the face of rapid environmental change requires forecasts from models of ecological processes. However, forecasts from ecological models are often associated with high degrees of uncertainty, making it difficult for such forecasts to inform decision-making processes. To make progress toward the goal of reliable and informative ecological forecasts, we need to know from where forecast uncertainty arises. Such knowledge can guide investment in future research that will most improve forecast skill. There is a rich history of analytical expressions that partition the variance of future dynamics, but these expressions suffer from necessary assumptions such as linear dynamics and small-variance approximations that exclude interactions. Similarly, the earth systems modeling community has developed methods for paritioning uncertainty of model projections, but these operate at a different modeling grain than most of ecology. Building on these approaches, we develop a simulation-based approach for

quantifying and partitioning forecast uncertainty from Bayesian state-space models that overcomes the limitations of previous analytical approaches. Our approach is similar to an Analysis of Variance, where the total variance of a forecast is paritioned among its constituent parts, namely initial conditions uncertainty, parameter uncertainty, driver uncertainty, process error, and their interactions. We apply the approach to near-term forecasts of the Yellowstone bison population and measles in China, demonstrating the broad utility of our approach. We also provide functions written in the statistical programming language R, which will allow others using Bayesian state-space models to employ our approach in their own research.

Introduction

Ecology is entering an era of prediction. This new era is being made possible by continuous data streams (e.g., the National Ecological Observatory Network, citizen science data, remote sensing), an increased statistical literacy among ecologists, and the need to provide actionable information for decision makers (Dietze et al. 2018). Thus, we are on the verge of answering Clark et al.'s (2001) challenge to make forecasting a core goal of ecology. Ecologists are in an excellent position to meet this forecasting challenge because we have spent decades gaining understanding of the processes that regulate populations, communities, and ecosystems. But we lack a systematic understanding of the current limits to ecological forecasts. As a result, we do not know how to allocate research effort to improve our forecasts.

Making poor forecasts is inevitable as ecology matures into a more predictive science. The key is to learn from our failures so that forecasts become more accurate over time. The success of meteorological forecasting tells us that basic research on the causes of forecast uncertainty is an essential component of this learning process (Bauer et al. 2015). Therefore, we need a systematic

44 and robust way to quantify and partition forecast uncertainty into its constituent parts (Dietze 2017).

45 Various approaches have been used to characterize and partition forecast uncertainty (Sobol'

46 1993, Cariboni et al. 2007). For example, consider a dynamic model designed to predict some state

47 y in the future (y_{t+1}) based on the current state (y_t), an environmental driver(s) (x), parameters (θ),

48 and process error (ε). We can then write a general form of the model as:

$$y_{t+1} = f(y_t, x_t | \theta) + \varepsilon_{t+1}, \tag{1}$$

49 which states that y at time $t+1$ is a function of y and x at time t conditional on the model

50 parameters (θ) plus process error (ε). Ignoring covariance among factors and assuming linear

51 dynamics, Dietze (2017), following Sobol' (1993) and Cariboni et al. (2007), suggests that forecast

52 variance ($Var[y_{t+1}]$) is approximately:

$$Var[y_{t+1}] \approx \underbrace{\left(\frac{\delta f}{\delta y}\right)^2}_{\text{stability}} \underbrace{Var[y_t]}_{\text{IC uncert.}} + \underbrace{\left(\frac{\delta f}{\delta x}\right)^2}_{\text{driver sens.}} \underbrace{Var[x_t]}_{\text{driver uncert.}} + \underbrace{\left(\frac{\delta f}{\delta \theta}\right)^2}_{\text{param sens.}} \underbrace{Var[\theta]}_{\text{param. uncert.}} + \underbrace{Var[\varepsilon_{t+1}]}_{\text{process error}}, \tag{2}$$

53 where each additive term follows a pattern of *sensitivity* times *variance* and "IC uncert." refers to

54 "*I*nitial *C*onditions uncertainty." The variance attributable to any particular factor is a function of

55 how sensitive the model is to the factor and the variance of that factor. For example, the atmosphere

56 is a chaotic system, meaning its dynamics are internally unstable and sensitive to initial conditions

57 uncertainty. This is why billions of dollars are spent each year to measure meterological variables –

58 meterologists learned that the key to reducing forecast error ($Var[y_{t+1}]$) was to reduce the

59 uncertainty of initial conditions ($Var[y_t]$). In contrast, ecologists are attempting to make actionable

60 forecasts with little knowledge of which term in Eq. 2 dominates forecast error. Knowing which

61 term dominates forecast error in different ecological settings will advance our fundamental

62 understanding of the natural world and immediately impact practical efforts to monitor, model, and

₆₃ predict ecological dynamics.

₆₄ While having an analytical expression such as Eq. 2 is satisfying, arriving at the expression

₆₅ involves strict assumptions. First, Eq. 2 only holds when the underlying dynamics are linear, which

₆₆ may not be the case for many populations and models. Second, Eq. 2 only includes additive effects

₆₇ of each factor because the Taylor series decomposition requires small-variance approximations that

₆₈ eliminate interactions. But, interactions among the factors are probably common. For example, in a

₆₉ simple simulation of an AR(1) process, we show that initial conditions uncertainty and parameter

₇₀ error interact to generate the full spread of forecast variance (Figure 1). Analyzing only the main

₇₁ effects of each source of uncertainty would lead to the false conclusion that initial conditions

₇₂ uncertainty is not important. Progress in quantifying and partioning forecast uncertainty therefore

₇₃ requires a more flexible approach than that provided by Eq. 2.

₇₄ We have four objectives in this paper. First, we introduce a canonical equation for error

₇₅ propagation and review approaches for quantifying sources of forecast uncertainty from other fields.

₇₆ Second, we describe a general, simulation-based method for quantifying and partitioning forecast

₇₇ uncertainty from Bayesian state-space models. Third, we apply our method to near-term forecasts

₇₈ of the Yellowstone bison population and measlse cases in China. These two applications

₇₉ demonstrate the generality of our approach. Fourth, through our applications we introduce

₈₀ functions written in the statistical programming language R (R Core Team 2016), which can be

₈₁ used to implement our method for many Bayesian state-space models.

Figure 1: Example of forecast uncertainty with different sources of error set to zero. Each line represents one realization, out of 200, from an order-one autoregressive model (AR(1) process). Contrary to the analytical expression (Eq. 2), initial conditions (I.C.) uncertainty and parameter uncertainty clearly interact. The spread of lines in (A) is not wholly because of initial conditions uncertainty or parameter uncertainty (compare panels B and C). It is their combined influence that causes the spread of realizations in (A and D). At least in this example, process error (set to zero in D) does appear to be independent, but we used a small value of process error to highlight other interactions. Source code: `generate_forecast_fxns.R`.

A Brief History of Quantifying and Partitioning Forecast Uncertainty

Error propagation

Methods for quantifying and partitioning forecast uncertainty all share common roots in statistical error propagation. Error propagation is concerned with translating the effects of variable (or parameter) uncertainty into the uncertainty of a function based those variables. That is, for the output q of some function $q = f(x_1, x_2, \ldots, x_n)$, we are interested in the variance of q given the variance of input variables ($\mathbf{x}$). It is this generic model formulation that leads to the canonical expression of error propagation

5

$$\sigma_q^2 = \left(\frac{\delta q}{\delta x_1}\sigma_{x_1}\right)^2 + \left(\frac{\delta q}{\delta x_2}\sigma_{x_2}\right)^2 + \cdots + \left(\frac{\delta q}{\delta x_n}\sigma_{x_n}\right)^2 \tag{3}$$

$$= \sum_{i=1}^{n}\left(\frac{\delta q}{\delta x_i}\sigma_{x_i}\right)^2. \tag{4}$$

This expression states that the variance of the output of function q (σ_q^2) is equal to the sum of squares of the input variables weighted by the sensitivity of the output variable to each input variable, quantified as the partial derivative. *Add more here on interactions, which require knowing the correlations among x's.*

Dynamic theoretical models

Lotka-Volterra style models and variance of *t+1* (draws from error propagation and sensitivity analysis).

Weather forecasting

Chaos and initial conditions.

Earth system models

Carbon and uncertainty from models and scenarios.

A Simulation-Based Approach for Partitioning Forecast Uncertainty

Analytical expressions of forecast uncertainty must rely on simplifying assumptions. Two important assumptions are (1) that different sources of uncertainty do not interact and (2) that the system of equations is linear. These analytical expressions are important for guiding our intuition, but these strict assumptions limit our ability to partition forecast uncertainty in practice. Thus, we present a simulation approach that is entirely model-based and requires no assumptions, other than those embedded in the model itself. We are building on the ideas put forth by Dietze (2017), who suggested a simulation approach for quantifying the terms in Eq. 2. Here we generalize the approach to fully partition forecast uncertainty among the main effects of different sources and their interactions.

As a starting point, consider a generic Bayesian state-space model

$$
\begin{aligned}
\textbf{Data Model:} \quad & y_t \sim \left[y_t \mid z_t, \sigma_{\mathrm{o}}^2 \right], && t = 1, \ldots, T, \\
\textbf{Process Models:} \quad & z_t \sim \left[z_t \mid \mu_t, \sigma_{\mathrm{p}}^2 \right], \\
& \mu_t = g\left(z_{t-1}, \mathbf{x}_t', \theta \right), && t = 2, \ldots, T, \\
\textbf{Parameter Models:} \quad & \phi \sim \left[\theta, \sigma_{\mathrm{p}}^2, \sigma_{\mathrm{o}}^2, z_{t=1} \right], && (5)
\end{aligned}
$$

where y_t is the observed state at time t, z_t is the latent state at time t, μ_t is the determinstic prediction of z at time t from the process model g, which is a function of z at time t-1, a vector of covariates ($\mathbf{x}$) at time t, and a set of unknown parameters, θ. σ_{o}^2 is observation error and σ_{p}^2 is process error. The notation $[a \mid b, c]$ reads, "the probability of a given b and c" (Gelfand and Smith 1990), and ϕ refers to the prior probability distributions for all unknown parameters and the initial conditions for the latent state, $z_{t=1}$.

For our purposes, we are interested in the probability distributions of the true state $\mathbf{z}$ at

future points in time, conditional on previous observations (**y**). This is referred to as the forecast distribution or the predictive process distribution (Hobbs and Hooten 2015 pp. 199–200), which, for one time step ahead of the final observation ($T + 1$), is defined as

$$
\begin{aligned}
[z_{T+1}|y_1,\ldots,y_T] = \int \int \ldots \int \left[z_{T+1}|z_T, \mathbf{x}_T, \theta, \sigma_p^2\right] \\
\times \left[z_1,\ldots,z_{T+1}, \theta, \sigma_p^2|y_1,\ldots,y_T\right] d\theta d\sigma_p^2 dz_1 \ldots dz_T.
\end{aligned}
\tag{6}
$$

The model in Eq. 5 can be fit using Markov chain Monte carlo (MCMC) algorithms, which makes calculating the forecast distribution a relatively simple task. To obtain $[z_{T+1}|y_1,\ldots,y_T]$, we can change the indexing of t in Eq. 5 to $t = 2,\ldots,T+1$ and then sample $z_{T+1}^{(k)}$ from $\left[z_{T+1}|g(z_T^{(k)}, \mathbf{x}_{T+1}^{(j(k))}, \theta^{(k)}), \sigma_p^{2(k)}\right]$ given the current values for $\theta^{(k)}$ and $z_T^{(k)}$ on each $k = 1,\ldots,K$ iteration of the MCMC algorithm (Hobbs and Hooten 2015, Williams et al. 2018). Note that we index the external covariate vector **x** using j and k, where $j(k)$ is realization j of the covariate vector **x** associated with MCMC sample k. In some cases, the external covariate will have only one value, in which case all K MCMC samples will share the same **x**. In other cases, their may be a distribution of external covariate values associated with uncertainty from the forecast of **x**, resulting in n values for each x_{T+1}. When $n < K$, which we anticipate it often will be, then **x** can be sampled with replacement and a value can be assigned to each MCMC sample. Making forecasts further into the future than $T + 1$ requires extending $T + 1$ to $T + 2,\ldots,T+q$ and iteratively sampling $\left[z_{T+q}|g(z_{T+q-1}^{(k)}, \mathbf{x}_{T+q}^{(j(k))}, \theta^{(k)}), \sigma_p^{2(k)}\right]$ (Hobbs and Hooten 2015).

The forecast distribution (Eq. 6) has all of the quantities that contribute to forecast uncertainty by incorporating their uncertainty explicitly across the K MCMC iterations. For example, initial conditions uncertainty is incorporated because $z_{T+1}^{(k)}$ is a function of $z_T^{(k)}$, resulting in a total of K point forecasts for z that comprise the posterior distribution of z at each time t. If we wish to ignore initial conditions uncertainty (I.C. uncertainty), we can make all K point forecasts starting from the mean of z_T

Table 1: Sampling equations for generating forecast distributions at times $T + q$ (where T is the time of the last observation) across $k = 1, \ldots, K$ MCMC samples with only certain sources of uncertainty present.

Source of Uncertainty	Notation	Sampling Equation	
I.C. Uncertainty	$V^{(I)} = V^{(I,\overline{PA},\overline{D},\overline{PS})}$	$\left[z_{T+q} \mid g(z^{(k)}_{T+q-1}, \mathbf{x}^{(*)}_T, \boldsymbol{\theta}^{(*)}), 0 \right]$	
Parameter Uncertainty	$V^{(PA)} = V^{(\overline{I},PA,\overline{D},\overline{PS})}$	$\left[z_{T+q} \mid g(z^{(k)}_{T+q-1}, \mathbf{x}^{(*)}_T, \boldsymbol{\theta}^{(k)}), 0 \right],$	$q > 1$
		$\left[z_{T+q} \mid g(z^{(*)}_T, \mathbf{x}^{(*)}_T, \boldsymbol{\theta}^{(k)}), 0 \right],$	$q = 1$
Driver Uncertainty	$V^{(D)} = V^{(\overline{I},\overline{PA},D,\overline{PS})}$	$\left[z_{T+q} \mid g(z^{(k)}_{T+q-1}, \mathbf{x}^{(j(k))}_T, \boldsymbol{\theta}^{(*)}), 0 \right],$	$q > 1$
		$\left[z_{T+q} \mid g(z^{(*)}_T, \mathbf{x}^{(j(k))}_T, \boldsymbol{\theta}^{(*)}), 0 \right],$	$q = 1$
Process Uncertainty	$V^{(PS)} = V^{(\overline{I},\overline{PA},\overline{D},PS)}$	$\left[z_{T+q} \mid g(z^{(k)}_{T+q-1}, \mathbf{x}^{(*)}_T, \boldsymbol{\theta}^{(*)}), \sigma^{2(k)}_{\mathrm{p}} \right],$	$q > 1$
		$\left[z_{T+q} \mid g(z^{(*)}_T, \mathbf{x}^{(*)}_T, \boldsymbol{\theta}^{(*)}), \sigma^{2(k)}_{\mathrm{p}} \right],$	$q = 1$

Note: The notation $V^{(A,\overline{B},\overline{C})}$ means that uncertainty from A enters the forecast from its posterior distribution (as approximated from MCMC samples), while sources B and C are set to their means (as calculated across the MCMC samples). In the main text we leave out the averaged terms (those with overbars) to reduce clutter. Thus, $V^{(A)} = V^{(A,\overline{B},\overline{C})}$.

$$z^{(*)}_T = E(z_T | y_1, \ldots, y_T) \approx \frac{\sum_{k=1}^{K} z^{(k)}_T}{K}, \tag{7}$$

143 which we call $z^{(*)}_T$ (as opposed to $z^{(k)}_T$), while retaining the uncertainty around all other parameters.

144 Our iterative sampling to obtain the forecast distribution then becomes a conditional statement,

$$z_{T+q} \sim \left[z_{T+q} | g(z^{(k)}_{T+q-1}, \mathbf{x}^{(j(k))}_T, \boldsymbol{\theta}^{(k)}), \sigma^{2(k)}_{\mathrm{p}} \right], \quad q > 1$$

$$z_{T+q} \sim \left[z_{T+q} | g(z^{(*)}_T, \mathbf{x}^{(j(k))}_T, \boldsymbol{\theta}^{(k)}), \sigma^{2(k)}_{\mathrm{p}} \right], \quad q = 1. \tag{8}$$

145 We can extend the basic idea of setting subsets of parameters and states to their posterior means (or

146 medians, depending on the distribution) to make partitioned forecasts where only prescribed

147 sources of uncertainty contribute to forecast uncertainty (Table 1).

148   It is important to note that although our discussion has centered on obtaining forecast

149   distributions within the MCMC algorithm used to fit the model, it is only feasible to do this for the

150   full forecast distribution (Eq. 6). In all other cases, where states or parameters must be averaged

151   over the K MCMC iterations, forecast simulations must be done *post hoc* using saved MCMC

152   samples (Box 1). In other words, estimating $z_1, z_2, \ldots, z_T$ is done within the MCMC fitting

153   algorithm, while estimating $z_{T+1}, z_{T+2}, \ldots, z_{T+q}$ is done after fitting the model, but with the full

154   MCMC output.

155   With these basics in mind, we now develop our approach for partitioning and quantifying

156   uncertainty, which is similar to an Analysis of Variance (ANOVA). Let the variance of the forecast

157   distribution at time $T + q$ be $V_{T+q}^{(X)}$, where $X = F$ (full forecast distribution), I (initial conditions

158   uncertainty only), PA (parameter uncertainty only), D (driver uncertainty), or PS (process

159   uncertainty only) (Table 1). $V_{T+q}^{(I)}$, $V_{T+q}^{(PA)}$, $V_{T+q}^{(D)}$, and $V_{T+q}^{(PS)}$ are the main effects of each factor on the

160   forecast distribution such that

$$
\begin{aligned}
V_{T+q}^{(F)} = {} & V_{T+q}^{(I)} + V_{T+q}^{(PA)} + V_{T+q}^{(D)} + V_{T+q}^{(PS)} \\
& + \varepsilon_{T+q}^{(I,PA)} + \varepsilon_{T+q}^{(I,D)} + \varepsilon_{T+q}^{(I,PS)} + \varepsilon_{T+q}^{(PA,PS)} + \varepsilon_{T+q}^{(PA,D)} + \varepsilon_{T+q}^{(D,PS)} \\
& + \varepsilon_{T+q}^{(I,PA,D)} + \varepsilon_{T+q}^{(I,PA,PS)} + \varepsilon_{T+q}^{(I,D,PS)} + \varepsilon_{T+q}^{(PA,D,PS)} \\
& + \varepsilon_{T+q}^{(I,PA,D,PS)},
\end{aligned} \tag{9}
$$

161   where the notation $\varepsilon_{T+q}^{(X,Y)}$ represents the remaining interactive effect of X and Y on $V_{T+q}^{(F)}$ after

162   accounting for their main effects. For example, if the full forecast variance is a function of only

163   initial conditions uncertainty I and parameter uncertainty PA, then

$$
V_{T+q}^{(F)} = V_{T+q}^{(I)} + V_{T+q}^{(PA)} + \varepsilon_{T+q}^{(I,PA)}, \tag{10}
$$

164   which rearranges to

$$\varepsilon_{T+q}^{(I,PA)} = V_{T+q}^{(F)} - \left[V_{T+q}^{(I)} + V_{T+q}^{(PA)} \right]. \tag{11}$$

165 We show an example of applying Eqs. 10-11 in a hypothetical situation where forecast uncertainty

166 is determined by initial conditions uncertainty and parameter uncertainty alone in Figure 2.

167   The necessary terms for partitioning forecast variance can be obtained by calculating the

168 variance of the partitioned forecast distributions (equations in Table 1 and combinations thereof).

169 To take this one step further, and to reiterate the core idea, let $V(F)$ be a function of initial

170 conditions uncertainty I, parameter uncertainty PA, and driver uncertainty D. We can then write the

171 equation for forecast uncertainty and the derived interaction effects as

> There must be a way to write a concise form of these equations vai summations.

$$V_{T+q}^{(F)} = V_{T+q}^{(I)} + V_{T+q}^{(PA)} + V_{T+q}^{(D)} + \varepsilon_{T+q}^{(I,PA)} + \varepsilon_{T+q}^{(I,D)} + \varepsilon_{T+q}^{(D,PA)} + \varepsilon_{T+q}^{(I,PA,D)}, \qquad \text{where} \tag{12}$$

$$\varepsilon_{T+q}^{(I,PA)} = V_{T+q}^{(I+PA)} - \left[V_{T+q}^{(I)} + V_{T+q}^{(PA)} \right] \tag{13}$$

$$\varepsilon_{T+q}^{(I,D)} = V_{T+q}^{(I+D)} - \left[V_{T+q}^{(I)} + V_{T+q}^{(D)} \right] \tag{14}$$

$$\varepsilon_{T+q}^{(PA,D)} = V_{T+q}^{(PA+D)} - \left[V_{T+q}^{(PA)} + V_{T+q}^{(D)} \right] \tag{15}$$

$$\varepsilon_{T+q}^{(I,PA,D)} = V_{T+q}^{(F)} - \left[V_{T+q}^{(I)} + V_{T+q}^{(PA)} + V_{T+q}^{(D)} \right] - \left[\varepsilon_{T+q}^{(I,PA)} + \varepsilon_{T+q}^{(I,D)} + \varepsilon_{T+q}^{(PA,D)} \right], \tag{16}$$

172 where the notation $V_{T+q}^{(A+B)}$ is the forecast variance under scenario where both A and B are allowed

173 to contribute to forecast uncertainty (i.e., a combination of the equations in Table 1). We present the

174 equations for the full partition among I, PA, D, and PS in Appendix 1. In the next two sections, we

175 apply our approach (Box 1) to quantify and partition the uncertainty of near-term forecasts of the

176 Yellowstone bison population and measles cases in China.

Figure 2: Graphical example of partitioning forecast uncertainty into main effects (V) and interaction effects (ε). The grey shaded area shows the interaction effect ($\varepsilon^{(I,PA)}$) that must be accounted for to fully partition forecast uncertainty between initial conditions uncertainty ($V^{(I)}$) and parameter uncertainty ($V^{(PA)}$). Source code: `generate_forecast_fxns.R`.

Box 1. Pseudocode for quantifying and partitioning forecast uncertainty from a Bayesian state-space model.

1. Fit a Bayesian state-space model (i.e., Eq. 5) with data $y_1, \ldots, y_T$ and save the MCMC samples.
2. Forecast $z_{T+q}^{(k)}$ for all $k = 1, \ldots, K$ MCMC samples to generate the full forecast distribution following Eq. 6 (this can be done within the MCMC algorithm or *post hoc* with saved MCMC samples).
3. Forecast $z_{T+q}^{(k)}$ for all $k = 1, \ldots, K$ MCMC samples to generate the partitioned forecast distributions for each source of uncertainty, averaging quantities over the K MCMC samples as necessary (equations in Table 1 and combinations thereof).
4. For each forecast time q, calculate the variance of each forecast distribution from steps 2-3.
5. Partition forecast variance using Eq. X.

Box 1:

Application: Yellowstone Bison Population

Application: Measles in China

Discussion

Acknowledgements

References

Bauer, P., A. Thorpe, and G. Brunet. 2015. The quiet revolution of numerical weather prediction. Nature 525:47–55.

Cariboni, J., D. Gatelli, R. Liska, and A. Saltelli. 2007. The role of sensitivity analysis in ecological modelling. Ecological Modelling 203:167–182.

Clark, J. S., S. R. Carpenter, M. Barber, S. Collins, A. Dobson, J. A. Foley, D. M. Lodge, M. Pascual, R. Pielke, W. Pizer, C. Pringle, W. V. Reid, K. A. Rose, O. Sala, W. H. Schlesinger, D. H. Wall, and D. Wear. 2001. Ecological forecasts: an emerging imperative. Science 293:657–660.

Dietze, M. C. 2017. Prediction in ecology: A first-principles framework. Ecological Applications 27:2048–2060.

Dietze, M. C., A. Fox, L. M. Beck-Johnson, J. L. Betancourt, M. B. Hooten, C. S. Jarnevich, T. H. Keitt, M. A. Kenney, C. M. Laney, L. G. Larsen, H. W. Loescher, C. K. Lunch, B. C. Pijanowski, J. T. Randerson, E. K. Read, A. T. Tredennick, R. Vargas, K. C. Weathers, and E. P. White. 2018.

Iterative near-term ecological forecasting: Needs, opportunities, and challenges. Proceedings of the National Academy of Sciences 115:1424–1432.

Gelfand, A. E., and A. F. Smith. 1990. Sampling-based approaches to calculating marginal densities. Journal of the American Statistical Association 85:398–409.

Hobbs, N. T., and M. B. Hooten. 2015. Bayesian Models: A Statistical Primer for Ecologists. Princeton University Press, Princeton.

R Core Team. 2016. R: A language and environment for statistical computing.

Sobol', I. 1993. Sensitivity Estimates for Nonlinear Mathematical Models.

Williams, P. J., M. B. Hooten, J. N. Womble, G. G. Esslinger, and M. R. Bower. 2018. Monitoring dynamic spatio-temporal ecological processes optimally. Ecology 99:524–535.